

Kalyan Big Data Projects – Project 10

How To Stream REGEX Data Into Hive Using Apache Flume

Pre-Requisites of Flume + Hive Project:

hadoop-2.6.0 , flume-1.6.0 , hive-1.2.1 , java-1.7

NOTE: Make sure that install all the above components

Flume + Hive Project Download Links:

`hadoop-2.6.0.tar.gz` ==> [link](https://archive.apache.org/dist/hadoop/core/hadoop-2.6.0/hadoop-2.6.0.tar.gz)
(<https://archive.apache.org/dist/hadoop/core/hadoop-2.6.0/hadoop-2.6.0.tar.gz>)

`apache-flume-1.6.0-bin.tar.gz` ==> [link](https://archive.apache.org/dist/flume/1.6.0/apache-flume-1.6.0-bin.tar.gz)
(<https://archive.apache.org/dist/flume/1.6.0/apache-flume-1.6.0-bin.tar.gz>)

`apache-hive-1.2.1-bin.tar.gz` ==> [link](http://mirror.fibergrid.in/apache/hive/hive-1.2.1/apache-hive-1.2.1-bin.tar.gz)
(<http://mirror.fibergrid.in/apache/hive/hive-1.2.1/apache-hive-1.2.1-bin.tar.gz>)

`kalyan-bigdata-examples.jar` ==> [link](https://github.com/kalyanhadooptraining/kalyan-bigdata-realtime-projects/blob/master/kalyan/kalyan-bigdata-examples.jar)
(<https://github.com/kalyanhadooptraining/kalyan-bigdata-realtime-projects/blob/master/kalyan/kalyan-bigdata-examples.jar>)

`kalyan-flume-hive-sink-1.6.0.jar` ==> [link](https://github.com/kalyanhadooptraining/kalyan-bigdata-realtime-projects/blob/master/kalyan/kalyan-flume-hive-sink-1.6.0.jar)
(<https://github.com/kalyanhadooptraining/kalyan-bigdata-realtime-projects/blob/master/kalyan/kalyan-flume-hive-sink-1.6.0.jar>)

`kalyan-regex-hive-agent.conf` ==> [link](https://github.com/kalyanhadooptraining/kalyan-bigdata-realtime-projects/blob/master/flume/project10-hive-regex/kalyan-regex-hive-agent.conf)
(<https://github.com/kalyanhadooptraining/kalyan-bigdata-realtime-projects/blob/master/flume/project10-hive-regex/kalyan-regex-hive-agent.conf>)

Learnings of this Project:

- We will learn Flume Configurations and Commands
- Flume Agent
 1. Source (Exec Source)
 2. Channel (Memory Channel)
 3. Sink (Hive Sink)
- Major project in Real Time `Product Log Analysis`
 1. We are extracting the data from server logs
 2. This data will be useful to do analysis on product views
 3. Complex Data is the output format then REGEX is best solution
- We can use Hive to analyze this data

1. create "**kalyan-regex-hive-agent.conf**" file with below content

```
agent.sources = EXEC
agent.sinks = HIVE
agent.channels = MemChannel

agent.sources.EXEC.type = exec
agent.sources.EXEC.command = tail -F /tmp/users.csv
agent.sources.EXEC.channels = MemChannel

agent.sinks.HIVE.type = hive
agent.sinks.HIVE.hive.metastore = thrift://localhost:9083
agent.sinks.HIVE.hive.database = kalyan
agent.sinks.HIVE.hive.table = users3
agent.sinks.HIVE.serializer = REGEX
agent.sinks.HIVE.serializer.regex = ^([^\,]*),([^\,]*),([^\,]*),([^\,]*),([^\,]*),([^\,]*),([^\,]*),([^\,]*)$
agent.sinks.HIVE.channel = MemChannel

agent.channels.MemChannel.type = memory
agent.channels.MemChannel.capacity = 1000
agent.channels.MemChannel.transactionCapacity = 100
```

2. Copy "**kalyan-regex-hive-agent.conf**" file into "**\$FUME_HOME/conf**" folder

3. Copy "**kalyan-bigdata-examples.jar** and **kalyan-flume-hive-sink-1.6.0.jar**" files into "**\$FLUME_HOME/lib**" folder

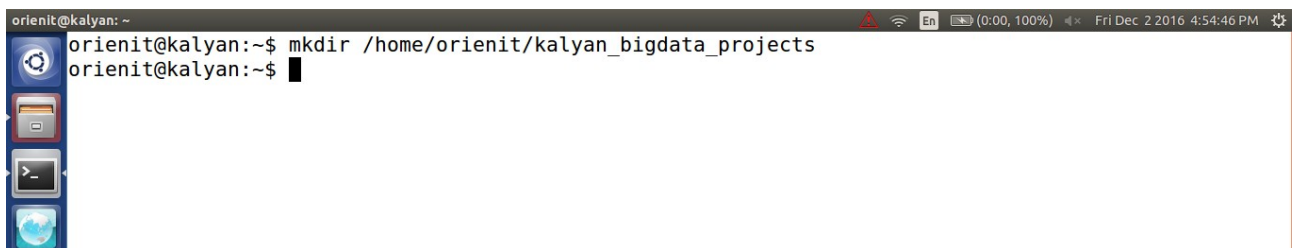
4. Generate Large Amount of Sample CSV data follow this [article](http://kalyanbigdatatraining.blogspot.com/2016/12/how-to-generate-large-amount-of-sample.html).

(<http://kalyanbigdatatraining.blogspot.com/2016/12/how-to-generate-large-amount-of-sample.html>)


5. Follow below steps...

i) Create '**kalyan_bigdata_projects**' folder in user home (i.e /home/orienit)

Command: `mkdir /home/orienit/kalyan_bigdata_projects`

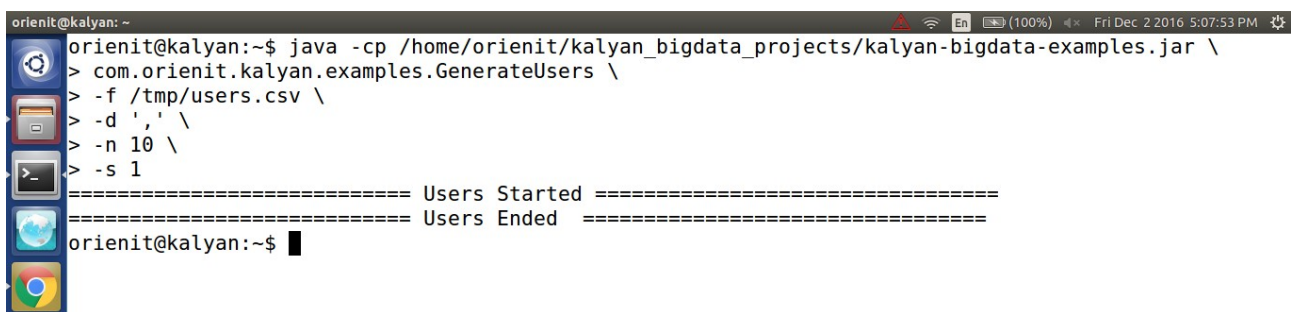


ii) Copy '**kalyan-bigdata-examples.jar**' jar file into '**/home/orienit/kalyan_bigdata_projects**' folder



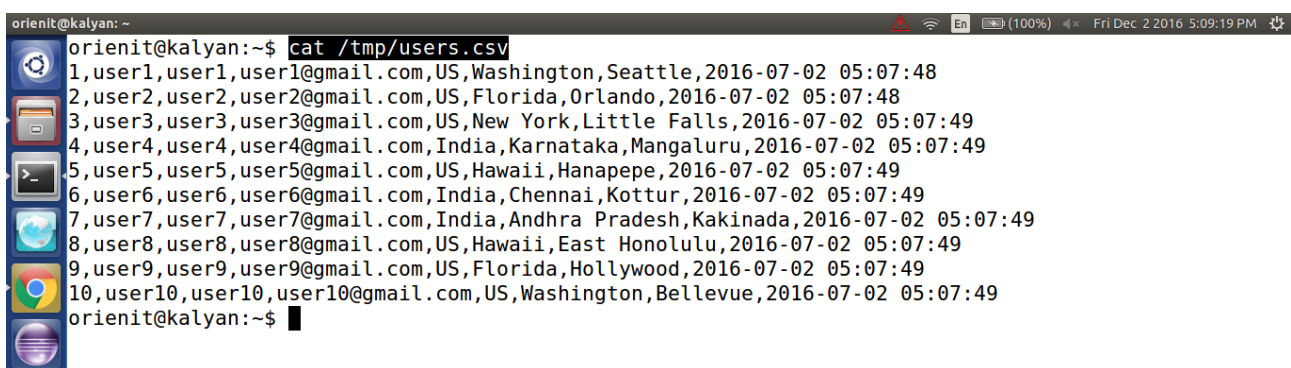
iii) Execute Below Command to Generate Sample CSV data with 100 lines. Increase this number to get more data ...

```
java -cp /home/orienit/kalyan_bigdata_projects/kalyan-bigdata-examples.jar \
com.orienit.kalyan.examples.GenerateUsers \
-f /tmp/users.csv \
-d ',' \
-n 100 \
-s 1
```



6. Verify the Sample CSV data in Console, using below command

```
cat /tmp/users.csv
```



7. To work with **Flume + Hive Integration**, Follow the below steps

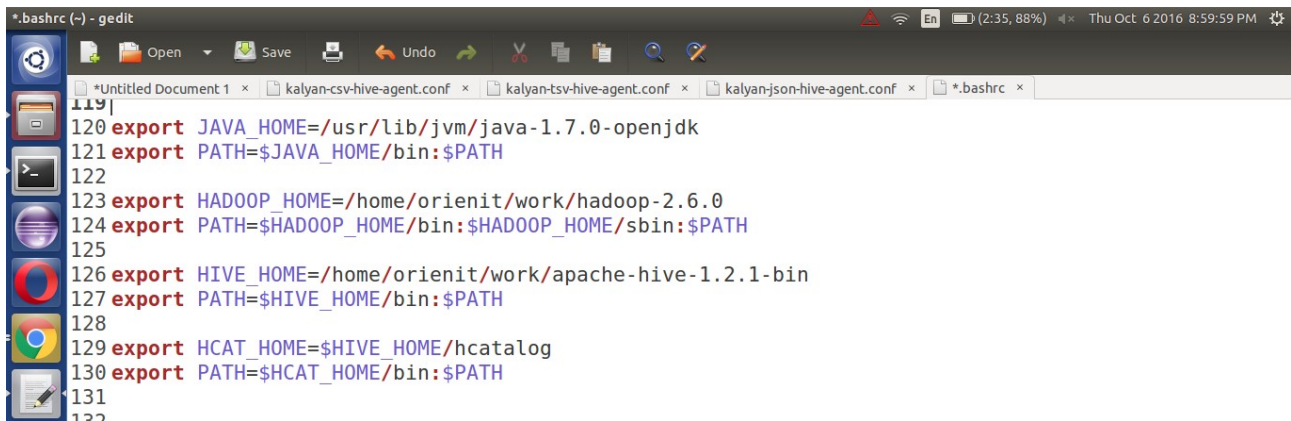
Follow this [article](#) to work with below procedure.

Refer: <http://kalyanbigdatatraining.blogspot.in/2016/10/how-to-work-with-acid-functionality-in.html>

i) update '~/.bashrc' file with below changes

```
export HIVE_HOME=/home/orienit/work/apache-hive-1.2.1-bin
export PATH=$HIVE_HOME/bin:$PATH
```

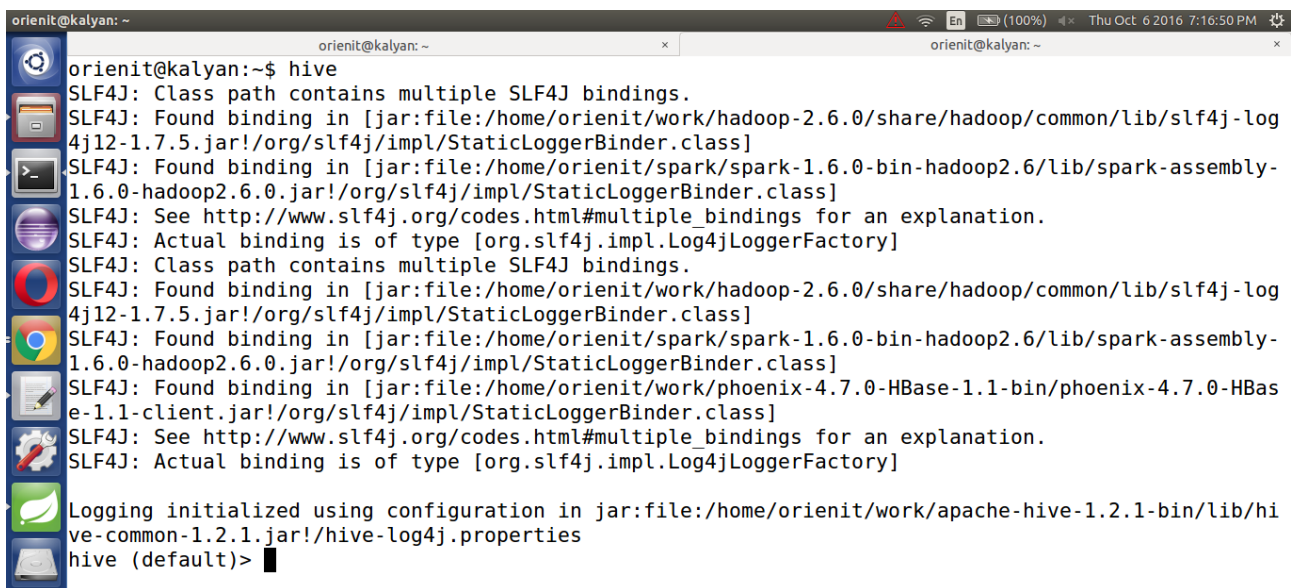
```
export HCAT_HOME=$HIVE_HOME/hcatalog
export PATH=$HCAT_HOME/bin:$PATH
```



```
*.bashrc (-) - gedit
119|
120 export JAVA_HOME=/usr/lib/jvm/java-1.7.0-openjdk
121 export PATH=$JAVA_HOME/bin:$PATH
122
123 export HADOOP_HOME=/home/orienit/work/hadoop-2.6.0
124 export PATH=$HADOOP_HOME/bin:$HADOOP_HOME/sbin:$PATH
125
126 export HIVE_HOME=/home/orienit/work/apache-hive-1.2.1-bin
127 export PATH=$HIVE_HOME/bin:$PATH
128
129 export HCAT_HOME=$HIVE_HOME/hcatalog
130 export PATH=$HCAT_HOME/bin:$PATH
131
132
```

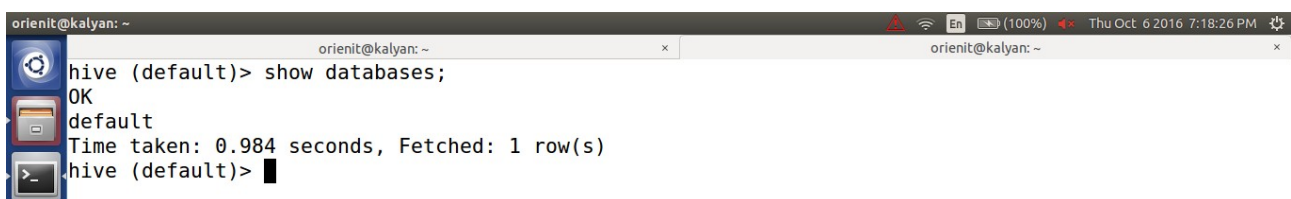
ii. reopen the Terminal

iii. start the hive using 'hive' command.



```
orienit@kalyan: ~
orienit@kalyan:~$ hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/orienit/work/hadoop-2.6.0/share/hadoop/common/lib/slf4j-log
4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/orienit/spark/spark-1.6.0-bin-hadoop2.6/lib/spark-assembly-
1.6.0-hadoop2.6.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/orienit/work/hadoop-2.6.0/share/hadoop/common/lib/slf4j-log
4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/orienit/spark/spark-1.6.0-bin-hadoop2.6/lib/spark-assembly-
1.6.0-hadoop2.6.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/orienit/work/phoenix-4.7.0-HBase-1.1-bin/phoenix-4.7.0-HBas
e-1.1-client.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
Logging initialized using configuration in jar:file:/home/orienit/work/apache-hive-1.2.1-bin/lib/hi
ve-common-1.2.1.jar!/hive-log4j.properties
hive (default)> █
```

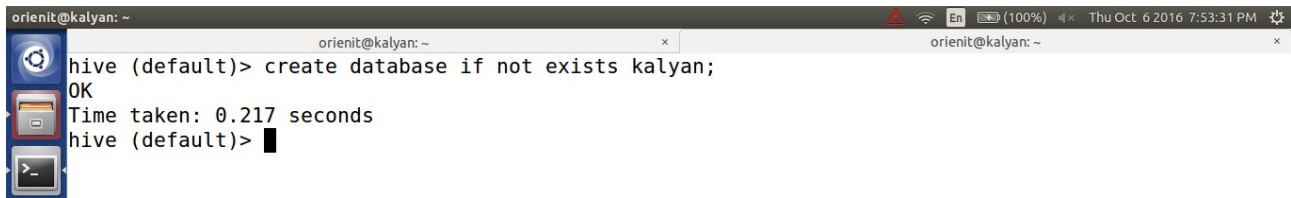
iv. list out all the databases in hive using 'show databases;' command



```
orienit@kalyan: ~
hive (default)> show databases;
OK
default
Time taken: 0.984 seconds, Fetched: 1 row(s)
hive (default)> █
```

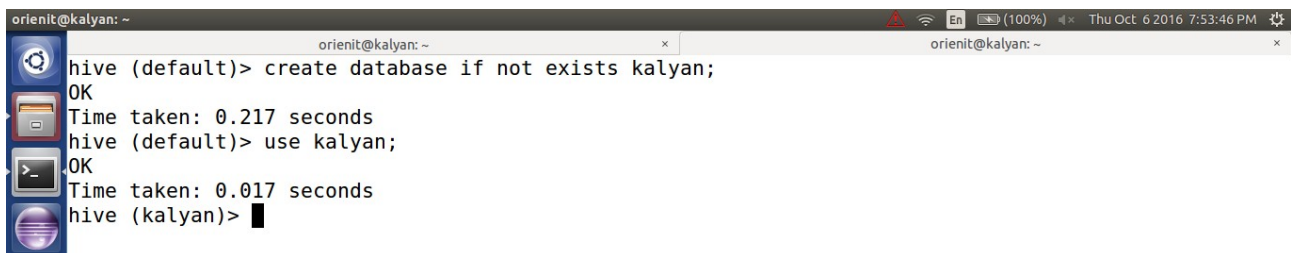
v. create a new database (**kalyan**) in hive using below command.

create database if not exists kalyan;



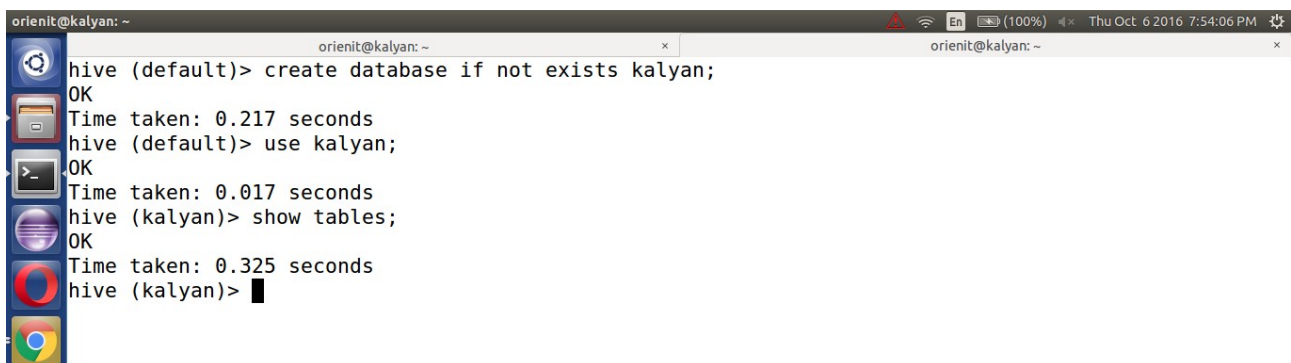
```
orientit@kalyan: ~  
orientit@kalyan: ~  
hive (default)> create database if not exists kalyan;  
OK  
Time taken: 0.217 seconds  
hive (default)> █
```

vi. use kalyan database using '**use kalyan;**' command



```
orientit@kalyan: ~  
orientit@kalyan: ~  
hive (default)> create database if not exists kalyan;  
OK  
Time taken: 0.217 seconds  
hive (default)> use kalyan;  
OK  
Time taken: 0.017 seconds  
hive (kalyan)> █
```

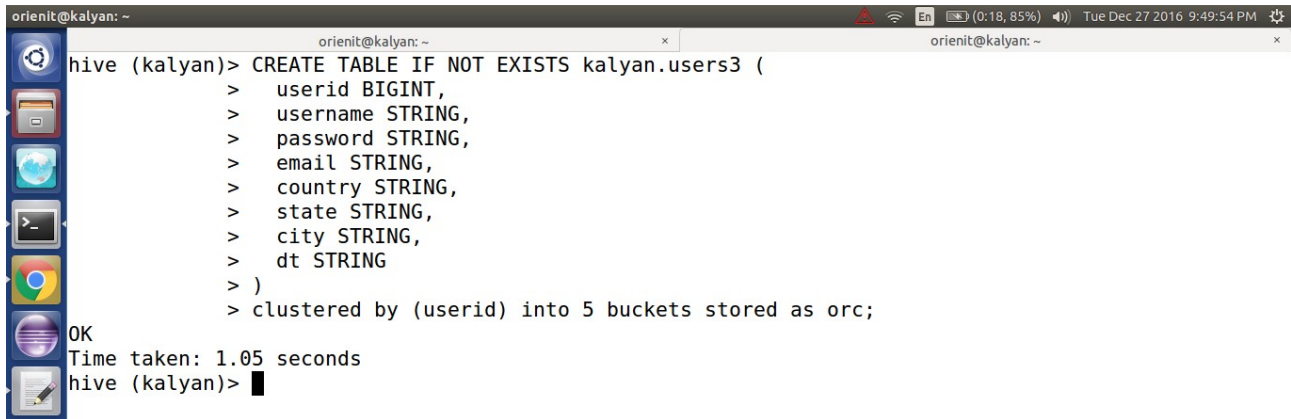
vii. list out all the tables in kalyan database using '**show tables;**' command.



```
orientit@kalyan: ~  
orientit@kalyan: ~  
hive (default)> create database if not exists kalyan;  
OK  
Time taken: 0.217 seconds  
hive (default)> use kalyan;  
OK  
Time taken: 0.017 seconds  
hive (kalyan)> show tables;  
OK  
Time taken: 0.325 seconds  
hive (kalyan)> █
```

viii. create '**users3**' table in kalyan database using below command.

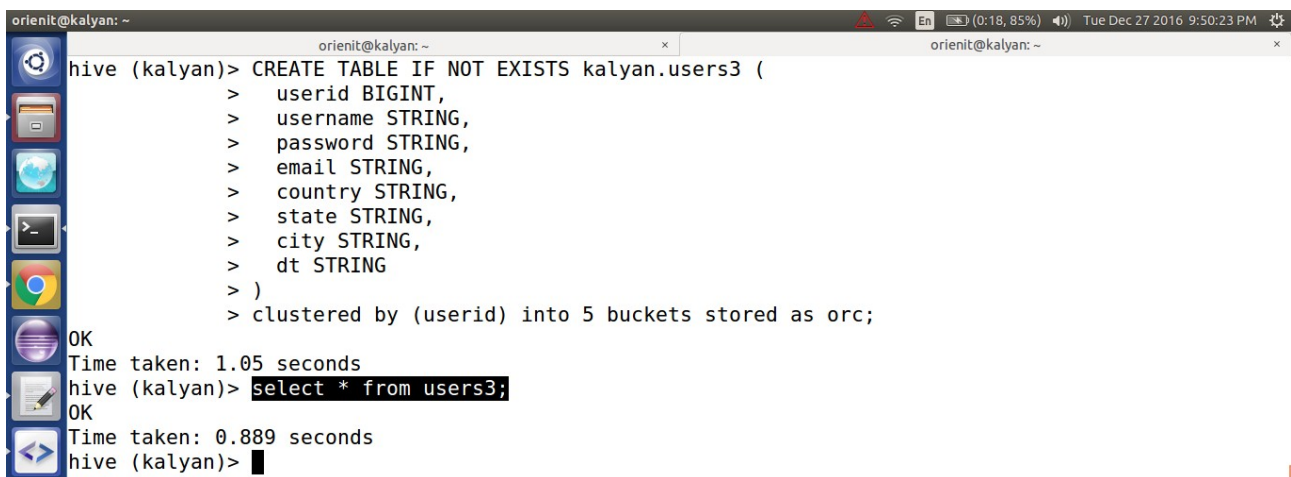
```
CREATE TABLE IF NOT EXISTS kalyan.users3 (  
  userid BIGINT,  
  username STRING,  
  password STRING,  
  email STRING,  
  country STRING,  
  state STRING,  
  city STRING,  
  dt STRING  
)  
clustered by (userid) into 5 buckets stored as orc;
```

```
orientit@kalyan: ~  
hive (kalyan)> CREATE TABLE IF NOT EXISTS kalyan.users3 (  
    >     userid BIGINT,  
    >     username STRING,  
    >     password STRING,  
    >     email STRING,  
    >     country STRING,  
    >     state STRING,  
    >     city STRING,  
    >     dt STRING  
    > )  
    > clustered by (userid) into 5 buckets stored as orc;  
OK  
Time taken: 1.05 seconds  
hive (kalyan)>
```

ix. Display the data from '**users3**' table using below command

`select * from users3;`



```
orientit@kalyan: ~  
hive (kalyan)> CREATE TABLE IF NOT EXISTS kalyan.users3 (  
    >     userid BIGINT,  
    >     username STRING,  
    >     password STRING,  
    >     email STRING,  
    >     country STRING,  
    >     state STRING,  
    >     city STRING,  
    >     dt STRING  
    > )  
    > clustered by (userid) into 5 buckets stored as orc;  
OK  
Time taken: 1.05 seconds  
hive (kalyan)> select * from users3;  
OK  
Time taken: 0.889 seconds  
hive (kalyan)>
```

x. start the hive in external metastore db mode using below command

`hive --service metastore`

```
orientit@kalyan: ~$ hive --service metastore
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/orientit/work/hadoop-2.6.0/share/hadoop/common/lib/slf4j-log
4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/orientit/spark/spark-1.6.0-bin-hadoop2.6/lib/spark-assembly-
1.6.0-hadoop2.6.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
Starting Hive Metastore Server
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/orientit/work/hadoop-2.6.0/share/hadoop/common/lib/slf4j-log
4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/orientit/spark/spark-1.6.0-bin-hadoop2.6/lib/spark-assembly-
1.6.0-hadoop2.6.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/orientit/work/phoenix-4.7.0-HBase-1.1-bin/phoenix-4.7.0-HBas
e-1.1-client.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
```

8. Execute the below command to **Extract data from REGEX data into Hive using Flume**

```
$FLUME_HOME/bin/flume-ng agent -n agent --conf $FLUME_HOME/conf -f
$FLUME_HOME/conf/kalyan-regex-hive-agent.conf -Dflume.root.logger=DEBUG,console
```

```
orientit@kalyan: ~$ $FLUME_HOME/bin/flume-ng agent -n agent --conf $FLUME_HOME/conf -f $FLUME_HOME/co
nf/kalyan-regex-hive-agent.conf -Dflume.root.logger=DEBUG,console
Info: Sourcing environment configuration script /home/orientit/work/apache-flume-1.6.0-bin/conf/flum
e-env.sh
Info: Including Hadoop libraries found via (/home/orientit/work/hadoop-2.6.0/bin/hadoop) for HDFS ac
cess
Info: Excluding /home/orientit/work/hadoop-2.6.0/share/hadoop/common/lib/slf4j-api-1.7.5.jar from cl
asspath
```

9. Verify the data in console

```
orientit@kalyan: ~$
ink.hive.HiveSink.drainOneBatch(HiveSink.java:299)] HIVE : Writing event to {metaStoreUri='thrift:/
/localhost:9083', database='kalyan', table='users1', partitionVals=[] }
2016-10-06 21:22:31,079 (SinkRunner-PollingRunner-DefaultSinkProcessor) [DEBUG - org.apache.flume.s
ink.hive.HiveSink.drainOneBatch(HiveSink.java:299)] HIVE : Writing event to {metaStoreUri='thrift:/
/localhost:9083', database='kalyan', table='users1', partitionVals=[] }
2016-10-06 21:22:31,079 (SinkRunner-PollingRunner-DefaultSinkProcessor) [DEBUG - org.apache.flume.s
ink.hive.HiveSink.drainOneBatch(HiveSink.java:299)] HIVE : Writing event to {metaStoreUri='thrift:/
/localhost:9083', database='kalyan', table='users1', partitionVals=[] }
```

10. Verify the data in Hive

Execute below command to get the data from hive table **'users3'**

```
select * from users3;
```

```

orienit@kalyan: ~
hive (kalyan)> select * from users3;
OK
1      user1      user1      user1@gmail.com India      Chennai Kottur      2016-52-27 09:52:13
2      user2      user2      user2@gmail.com US        Florida Orlando    2016-52-27 09:52:13
3      user3      user3      user3@gmail.com US        Florida Hollywood   2016-52-27 09:52:14
4      user4      user4      user4@gmail.com India      Karnataka Bengaluru   2016-52-27 09:52:14
5      user5      user5      user5@gmail.com India      Karnataka Bengaluru   2016-52-27 09:52:14
6      user6      user6      user6@gmail.com US        Florida Hollywood   2016-52-27 09:52:14
7      user7      user7      user7@gmail.com India      Telangana Nizamabad    2016-52-27 09:52:14
8      user8      user8      user8@gmail.com India      Andhra Pradesh Rajahmundry  2016-52-27 09:52:14
9      user9      user9      user9@gmail.com US        New York Canandaigua   2016-52-27 09:52:14
10     user10     user10     user10@gmail.com US        Hawaii Pearl City    2016-52-27 09:52:14
11     user11     user11     user11@gmail.com India      Andhra Pradesh Kakinada     2016-52-27
09:52:14
12     user12     user12     user12@gmail.com India      Chennai Kottur    2016-52-27 09:52:15
13     user13     user13     user13@gmail.com US        Washington Bellevue      2016-52-27
09:52:15
14     user14     user14     user14@gmail.com India      Andhra Pradesh Guntur      2016-52-27 09:52:15
15     user15     user15     user15@gmail.com India      Chennai Adambakkam 2016-52-27 09:52:15
16     user16     user16     user16@gmail.com US        Hawaii Mililani Town 2016-52-27 09:52:15
17     user17     user17     user17@gmail.com US        Washington Bellevue      2016-52-27
09:52:15

```